

CUDA Thread Hierarchy

Memory Scope

Register
per-thread


Local Memory
per-thread

Shared Memory
per-block

L1 / L2 Cache
on-chip

Global Memory
all threads

Constant Memory
SM read-only



1 Grid = many Blocks; 1 Block = many Warps; 1 Warp = 32 Threads; 1 Block runs on 1 SM