

Three complementary eval layers

Layer	Best for	Strength	Main risk
code eval	format, fields, regex, length	fast, cheap, reproducible	brittle on semantic quality
LLM-as-a-judge	correctness, tone, faithfulness	understands meaning	judge prompt can drift or be wrong
human evaluation	new failure modes, ground truth	domain judgment	slow, costly, fatigue-prone

A serious suite stacks layers because each one leaves different holes.