

Fair-failure inspection: CORE-Bench story

A failing trace should reveal a real mistake, not an eval artifact

Strict exact match

Reference: 96.12
Model: 96.124991
Result: FAIL



Numeric tolerance

$|96.124991 - 96.12| = 0.004991$
Within tolerance
Result: PASS

Reported score before eval repair



Reported score after eval repair



Speaker example; exact benchmark details should be treated as provenance-bound to the talk.

Source: 01:30:40--01:31:45