


Build evals in phases; rerun the full dataset at gates

Small sets are for learning; full datasets are for overfit and regression checks



Keep agent version, rubric version, dataset version, and judge version attached to every result.