


Parser/token economy pipeline

Let the model design a parser once, then let code process many pages.



Token intuition

LLM parses every page: Cost $\approx N \cdot H$
LLM builds parser once: Cost $\approx H + N \cdot s$
Assumption: $s \ll H$

The speaker claims about 99% token savings for this style of workflow.