

Retry amplification turns small errors into GPU spikes

$$\text{effective load} = \text{base demand} \times \text{retries} \times \text{tool cost}$$

